


docker

Day 4 and Day 5

Docker Swarm



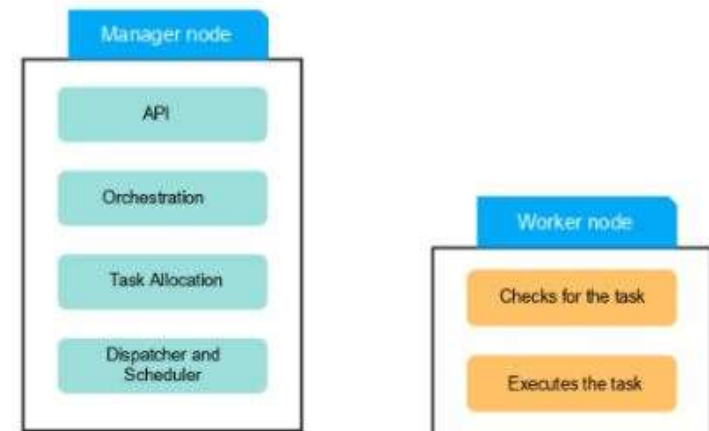
- Docker Swarm is an orchestration management tool that runs on Docker applications. It helps end-users in creating and deploying a cluster of Docker nodes.
- Each node of a Docker Swarm is a Docker daemon, and all Docker daemons interact using the Docker API. Each container within the Swarm can be deployed and accessed by nodes of the same cluster.

Key Concepts

- Service and Tasks
 - Docker containers are launched using services.
 - Services can be deployed in two different ways - global and replicated.
 - Global services are responsible for monitoring containers that want to run on a Swarm node. In contrast, replicated services specify the number of identical tasks that a developer requires on the host machine.
 - Services enable developers to scale their applications.
- Node
 - A Swarm node is an instance of the Docker engine.
 - It is possible to run multiple nodes on a single server. But in production deployments, nodes are distributed across various devices.

Key Concepts

- There are three types of nodes in Docker Swarm:
 - Manager node: Maintains cluster management tasks
 - Worker node: Receives and executes tasks from the manager node
 - Leader node: makes all of the swarm management and task orchestration decisions for the swarm

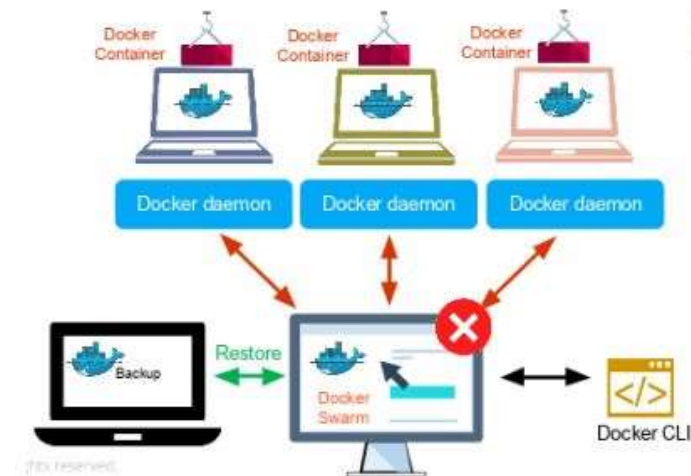


Working with Swarm

- A service is created based on the command-line interface.
- The API that we connect in our Swarm environment allows us to do orchestration by creating tasks for each service.
- The task allocation will enable us to allocate work to tasks via their IP address.
- The dispatcher and scheduler assign and instruct worker nodes to run a task.
- The Worker node connects to the manager node and checks for new tasks.
- The final stage is to execute the tasks that have been assigned from the manager node to the worker node.

Failure Correction Mechanism

- Consider an environment with docker containers
- If one of the containers fails, we can use the Swarm to correct that failure. Docker Swarm can reschedule containers on node failures. Swarm node has a backup folder which we can use to restore the data onto a new Swarm.

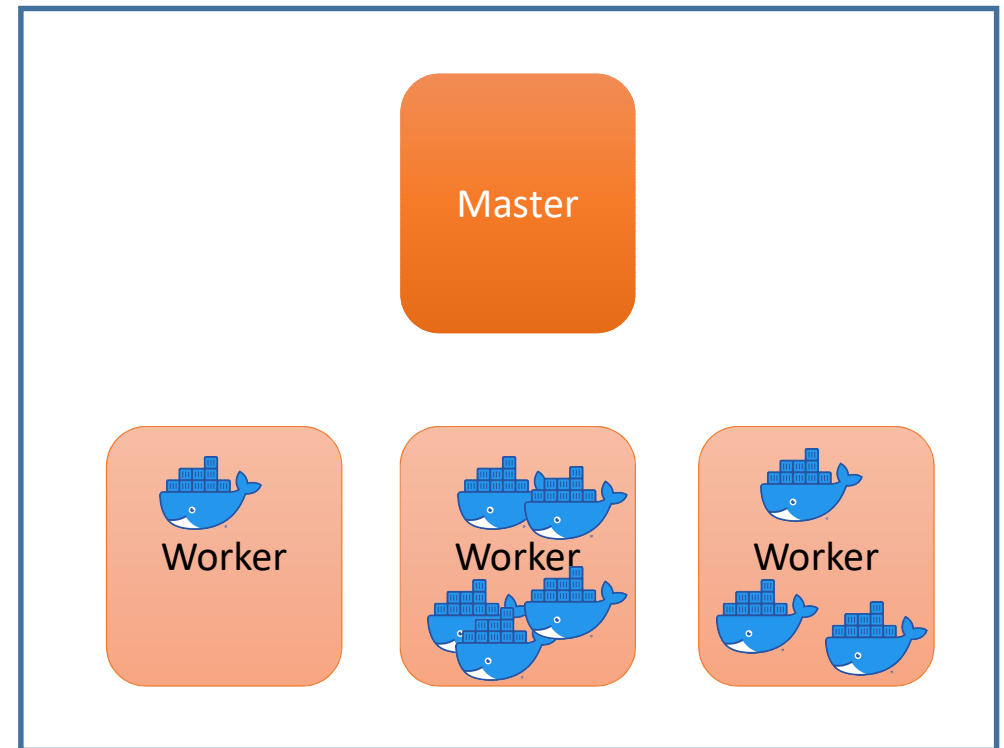


Kubernetes

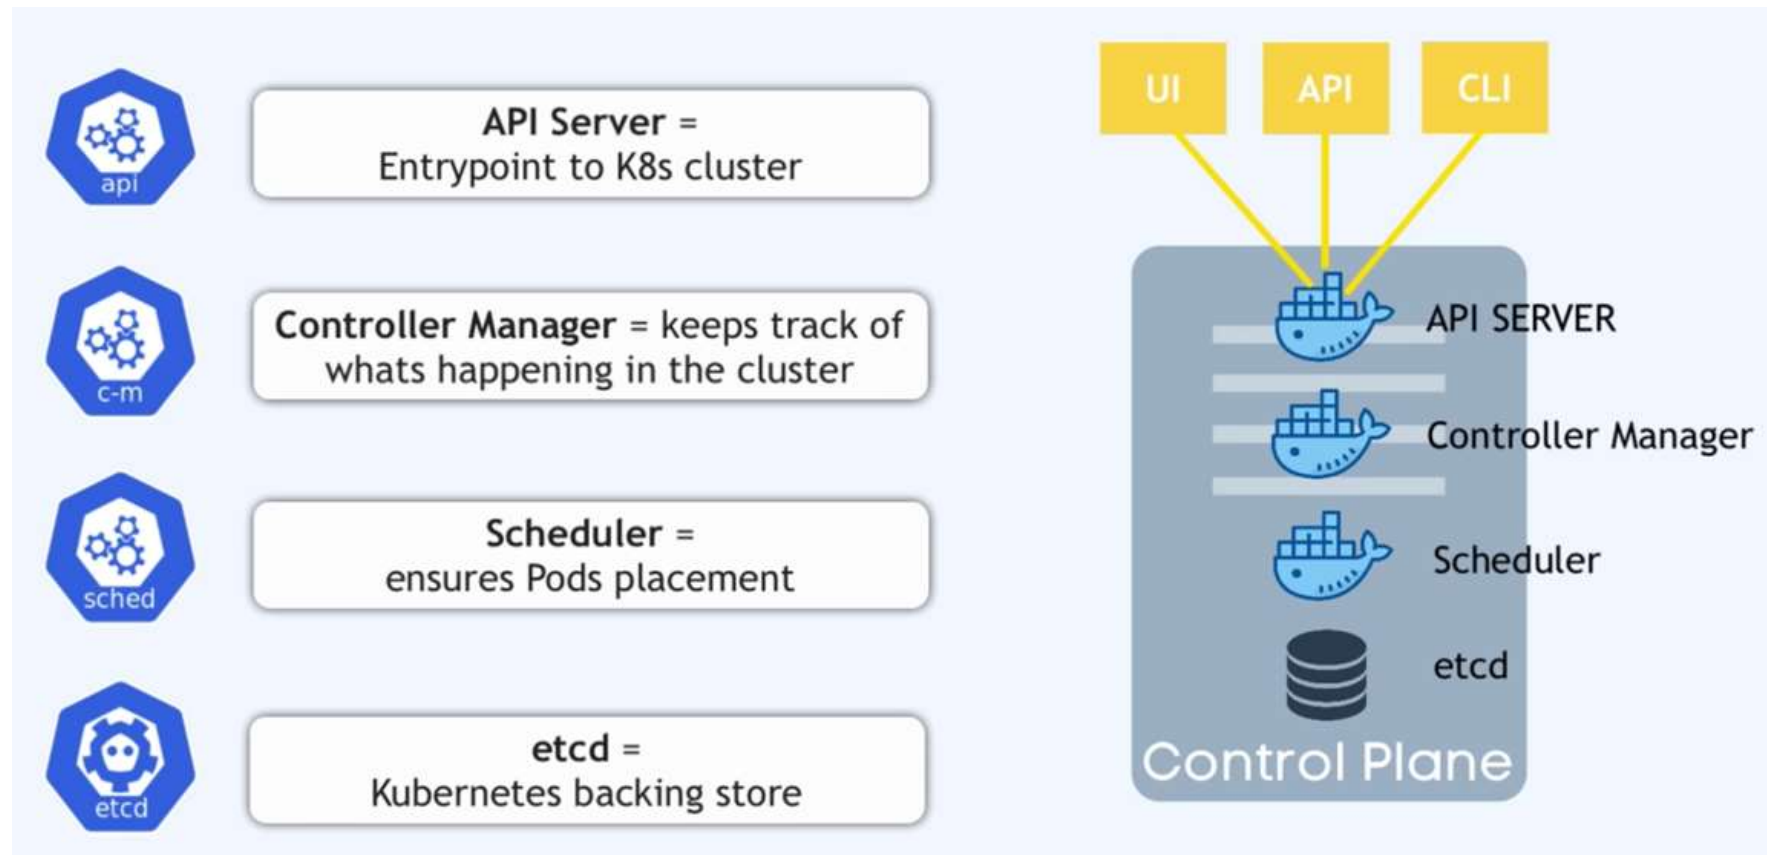
- Kubernetes is a container orchestration tool
- Helps manage containerized applications in different deployment environments such as physical machines, virtual machines, cloud and hybrid environments
- Need for orchestration tool:
 - Raise from monolith to microservices
 - Increases use of containers and need for a way to run them
 - Need for the management of containers

Kubernetes

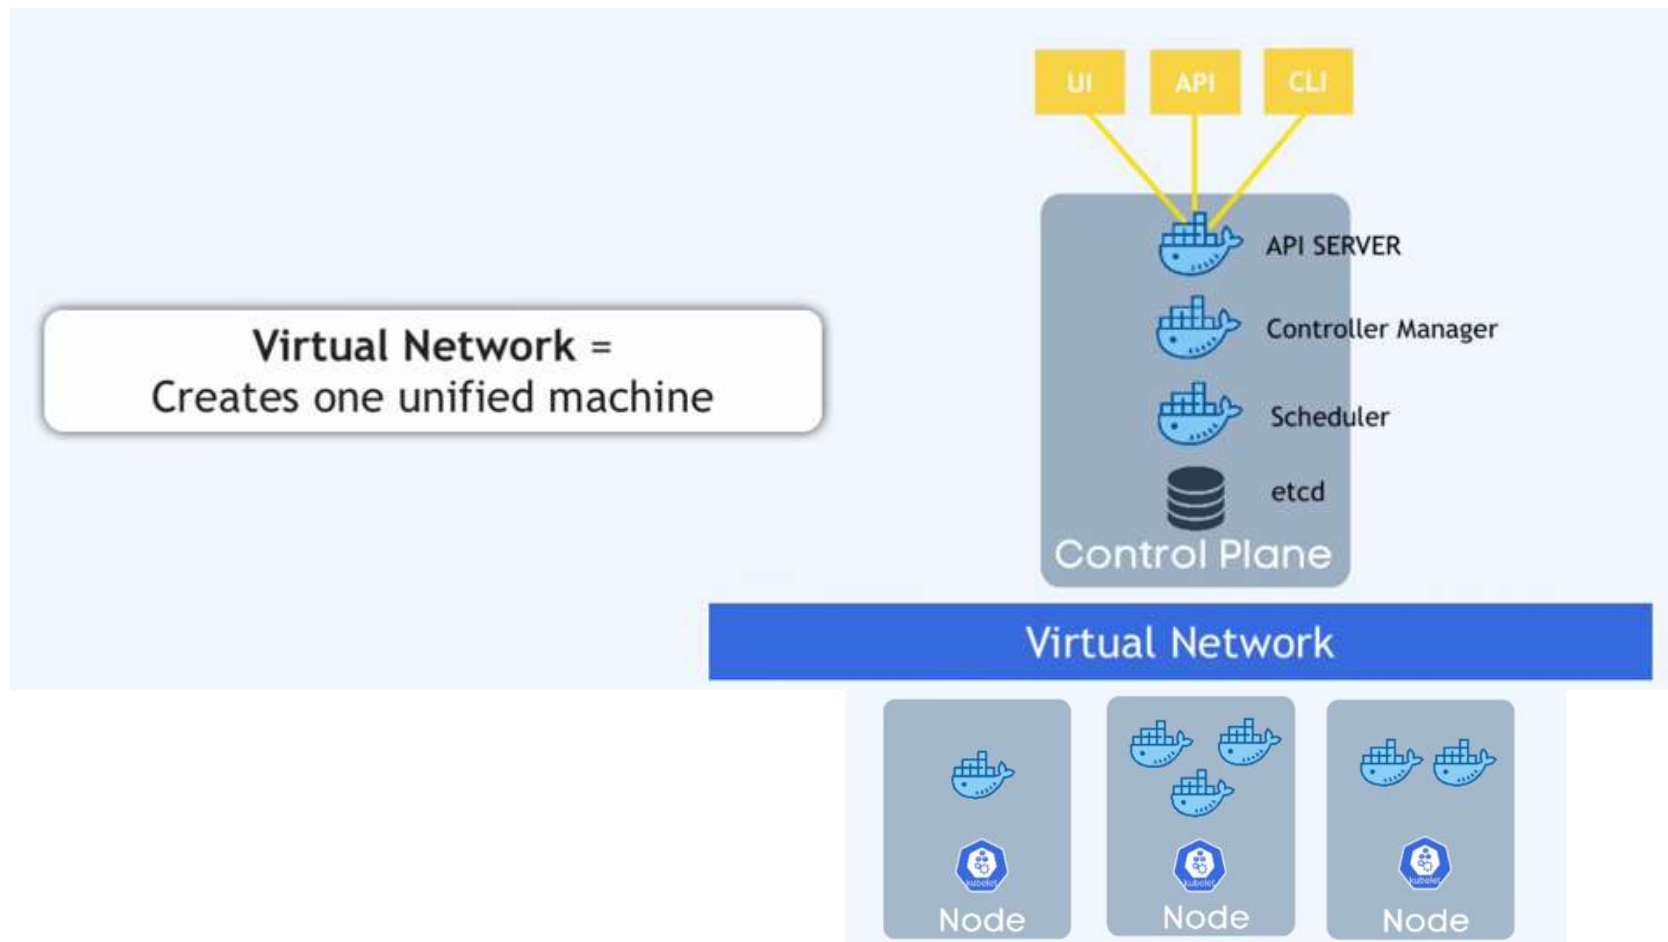
- Features of orchestration tools:
 - High availability and zero down time
 - Scalability or high performance
 - Disaster recovery – backup and restore
- Kubernetes Architecture:
 - Master Node
 - Worker Nodes



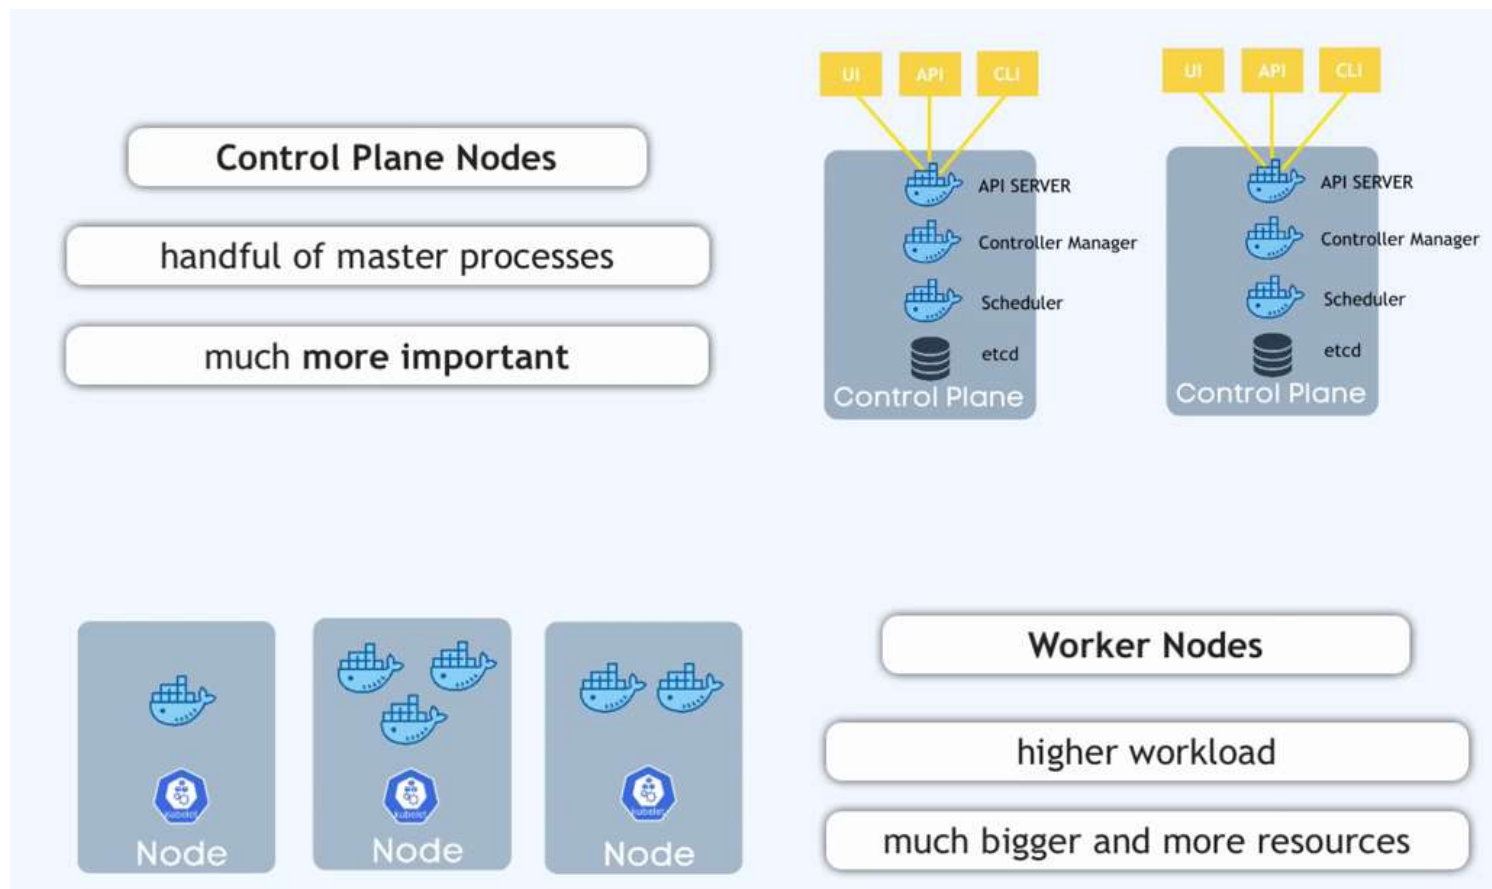
Master Node

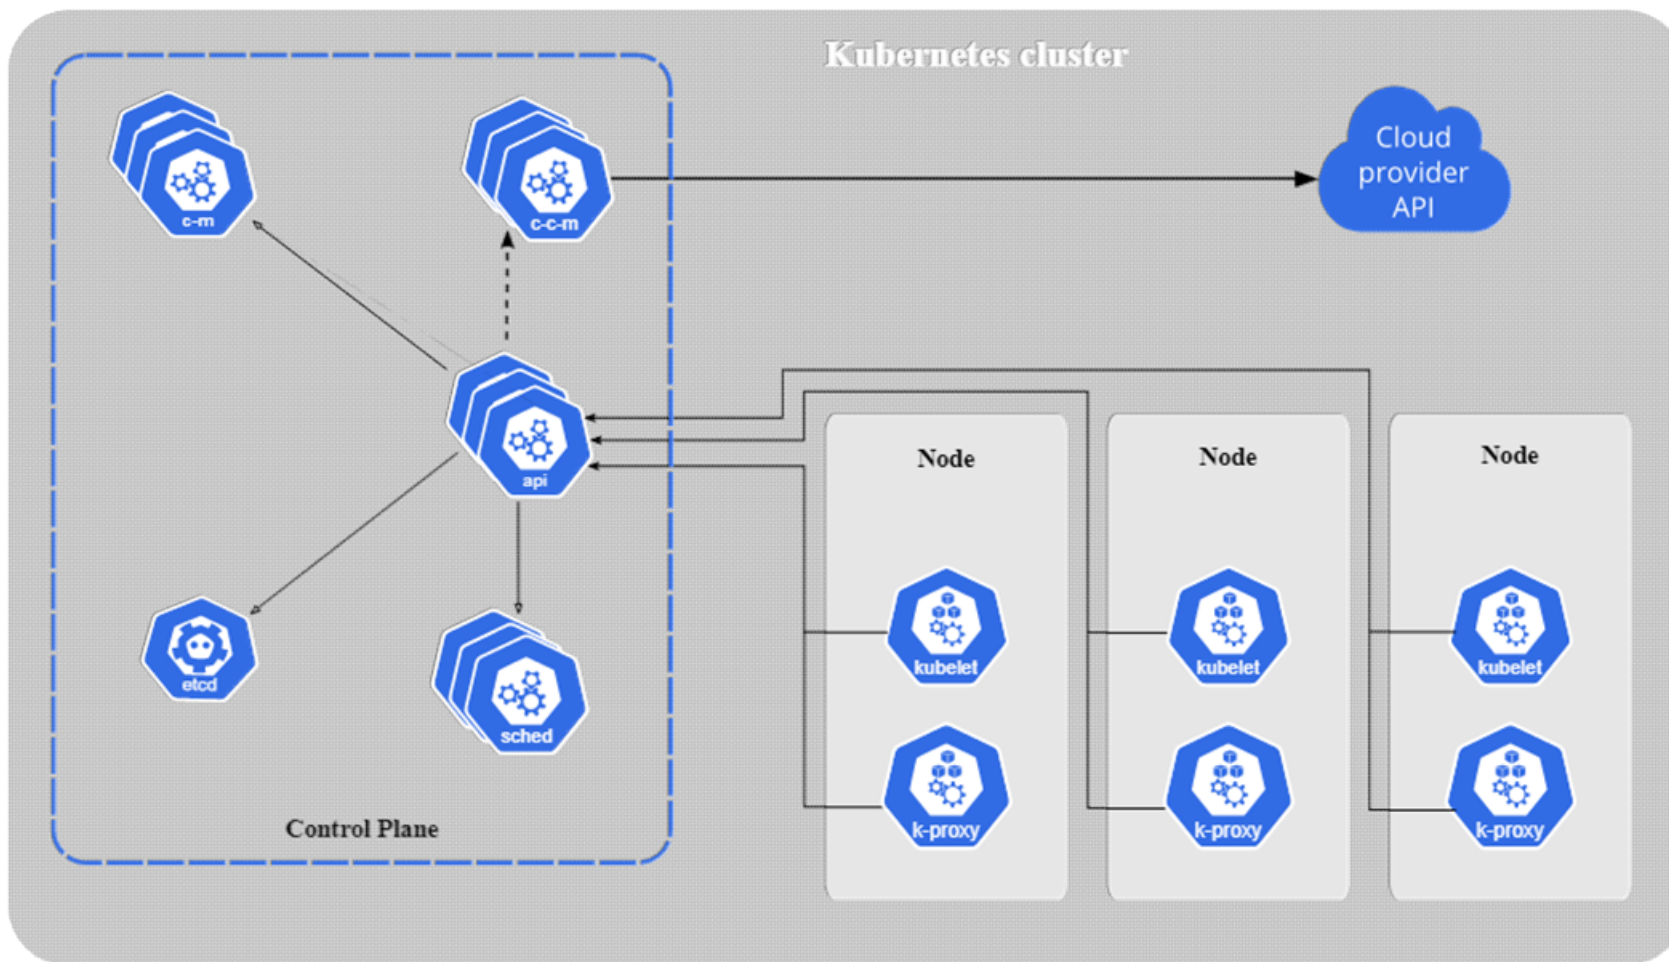


Virtual Network

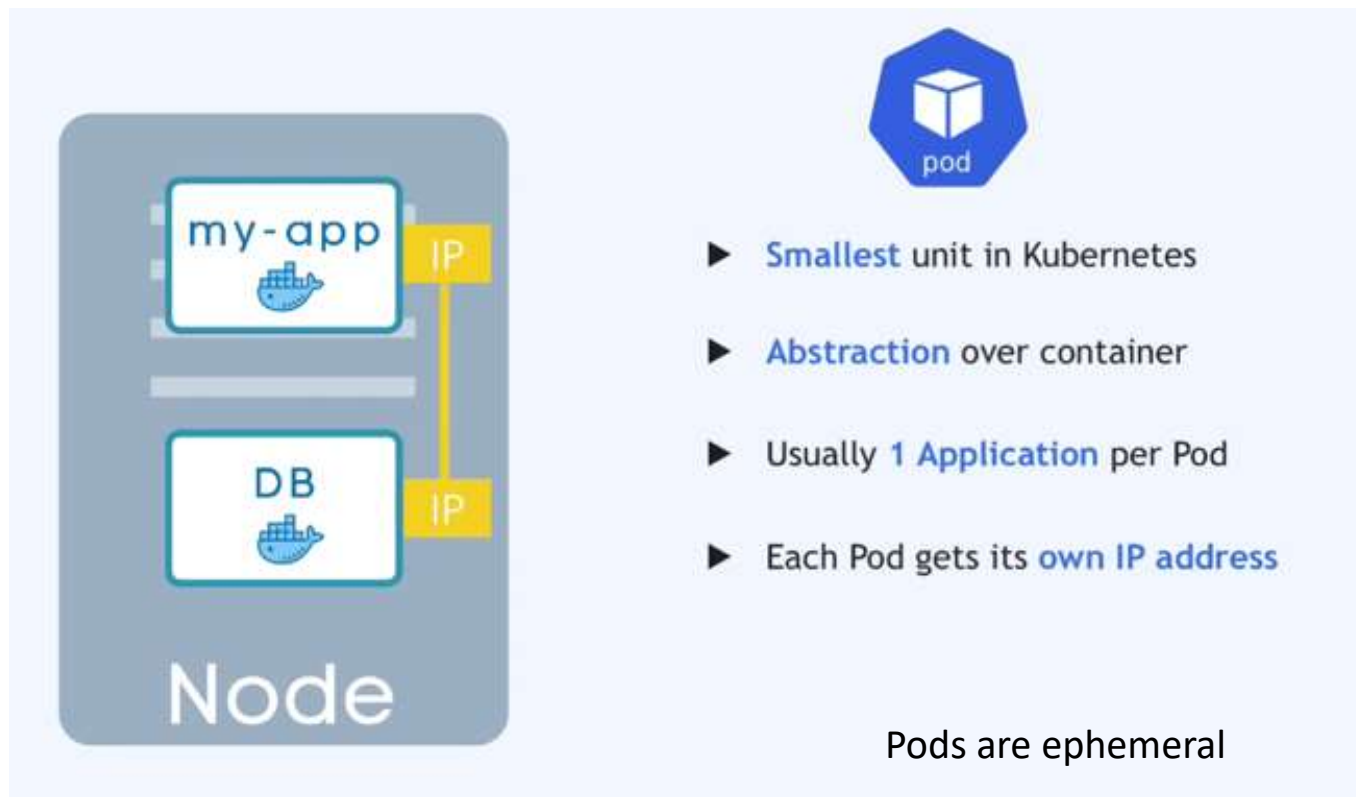


Control Plane Backups

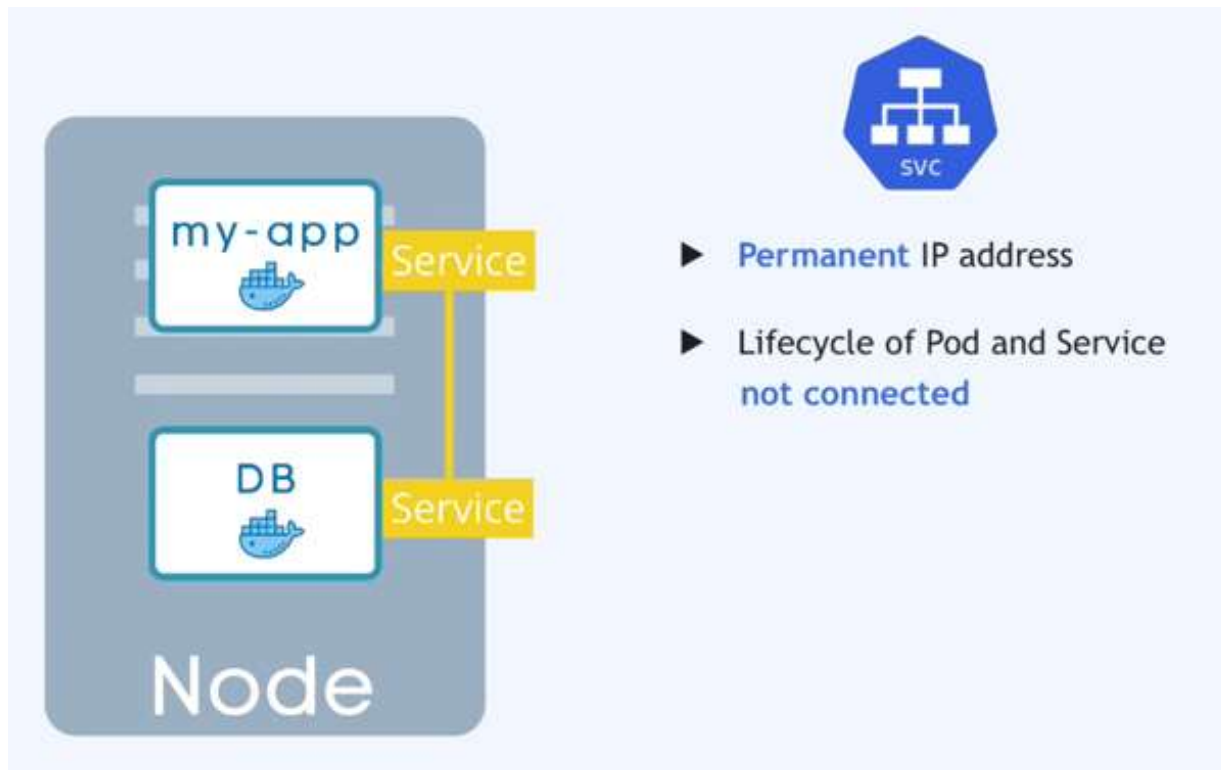




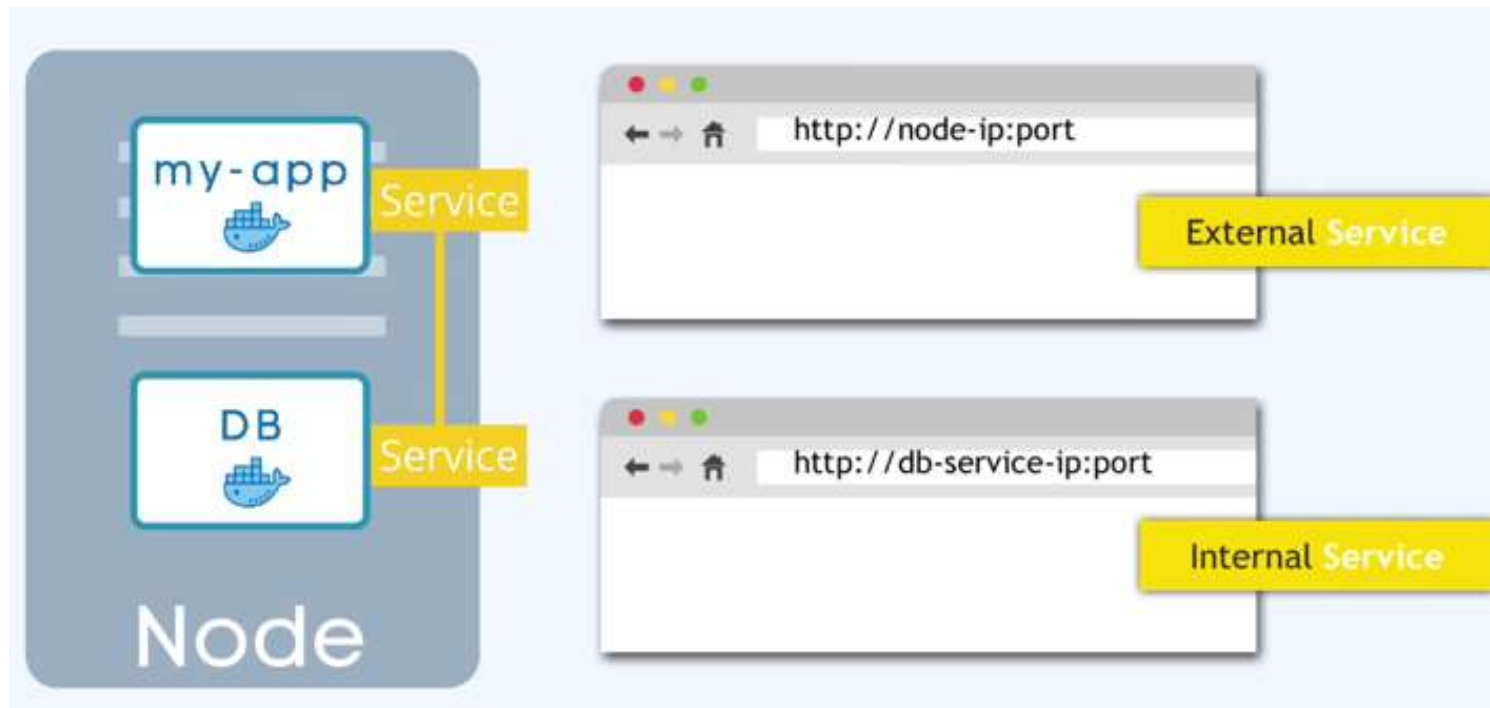
Nodes and Pods



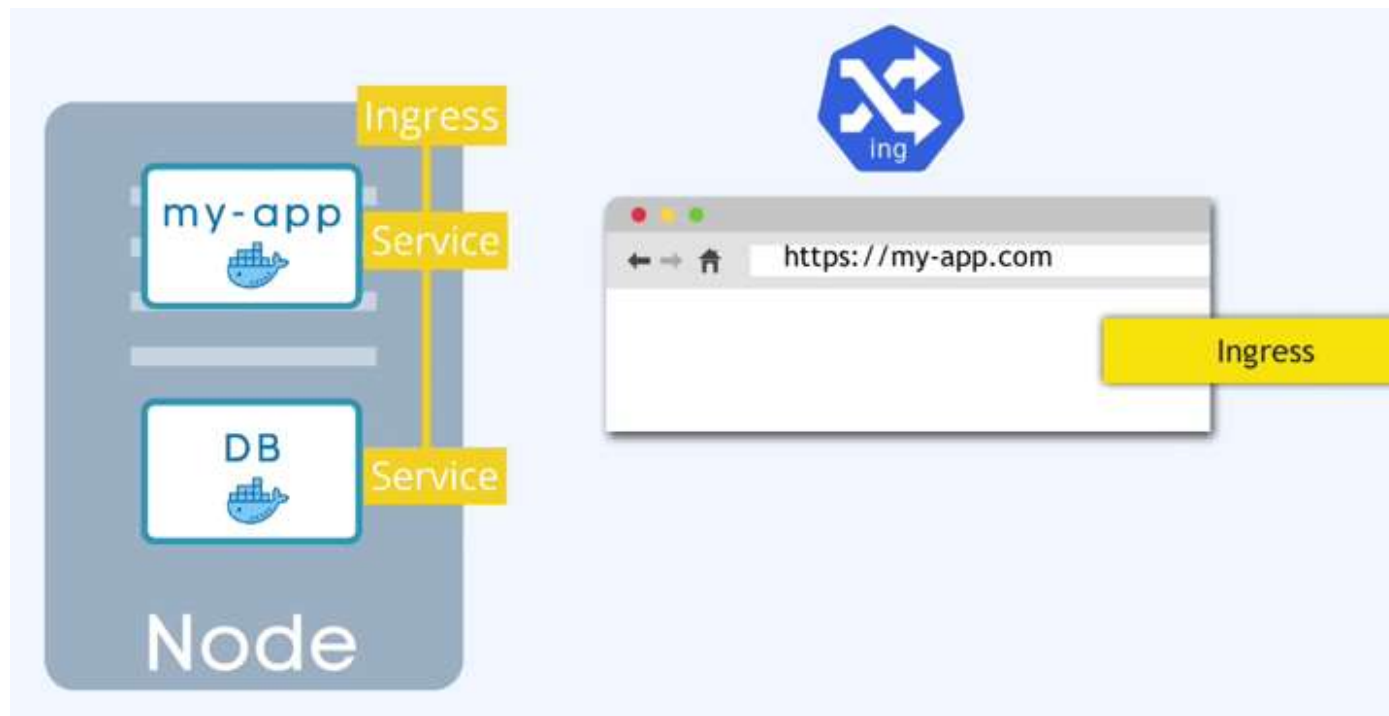
Service



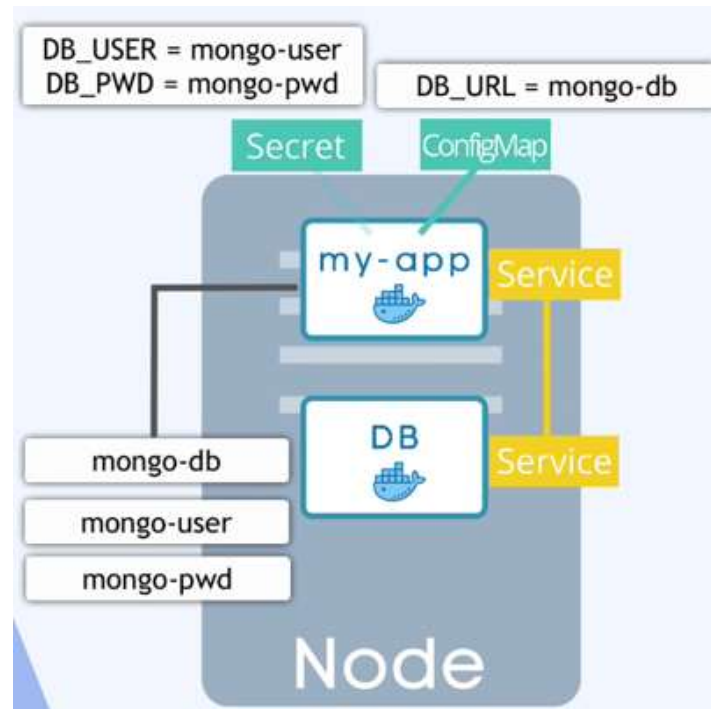
Service



Ingress

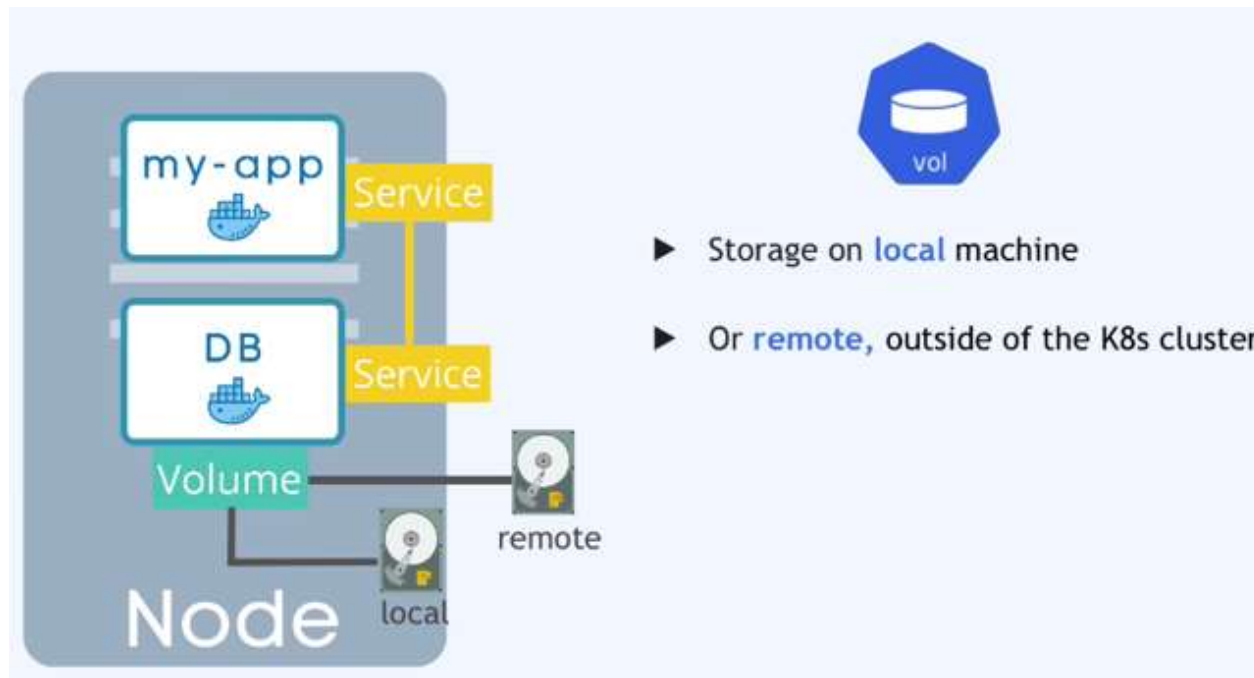


Configmap and Secret



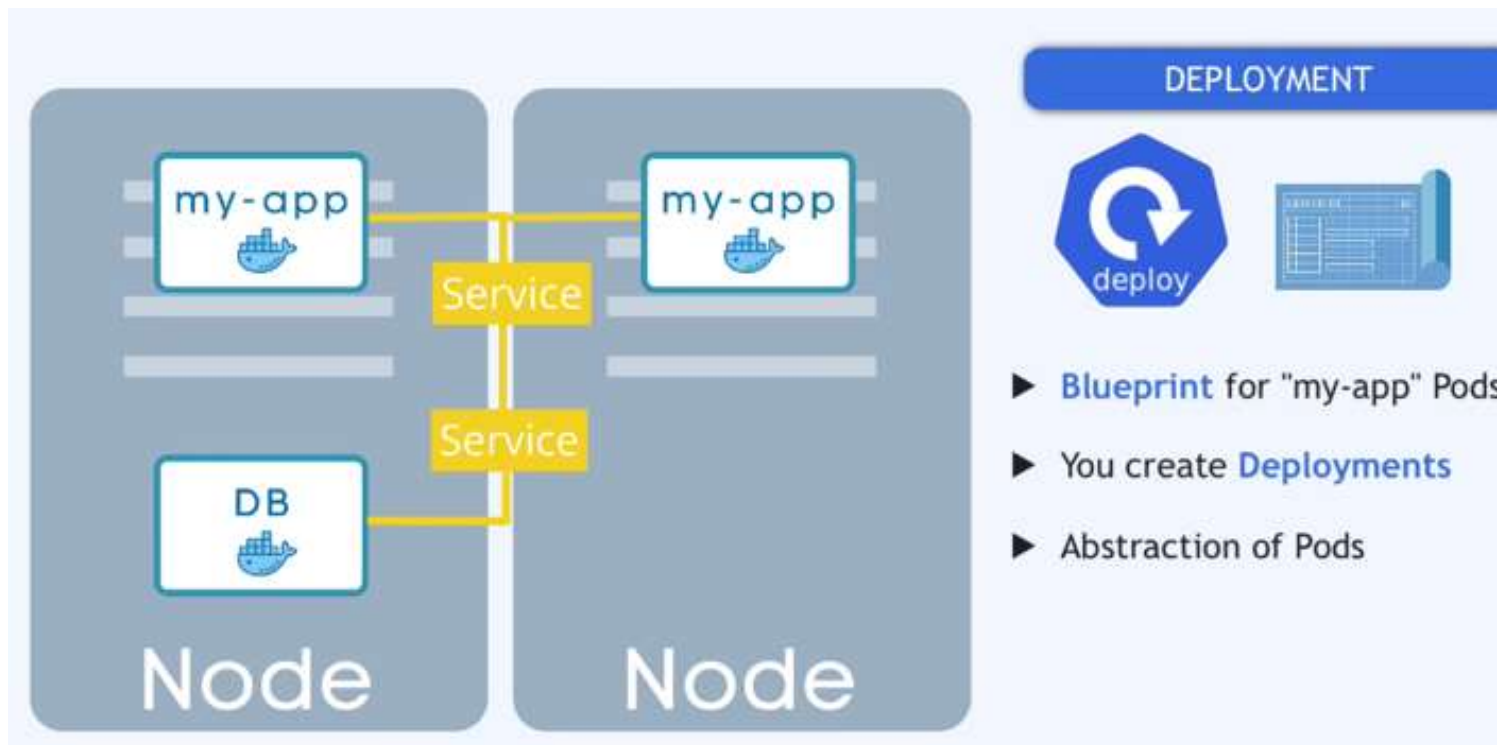
Reference them in the pod

Volumes



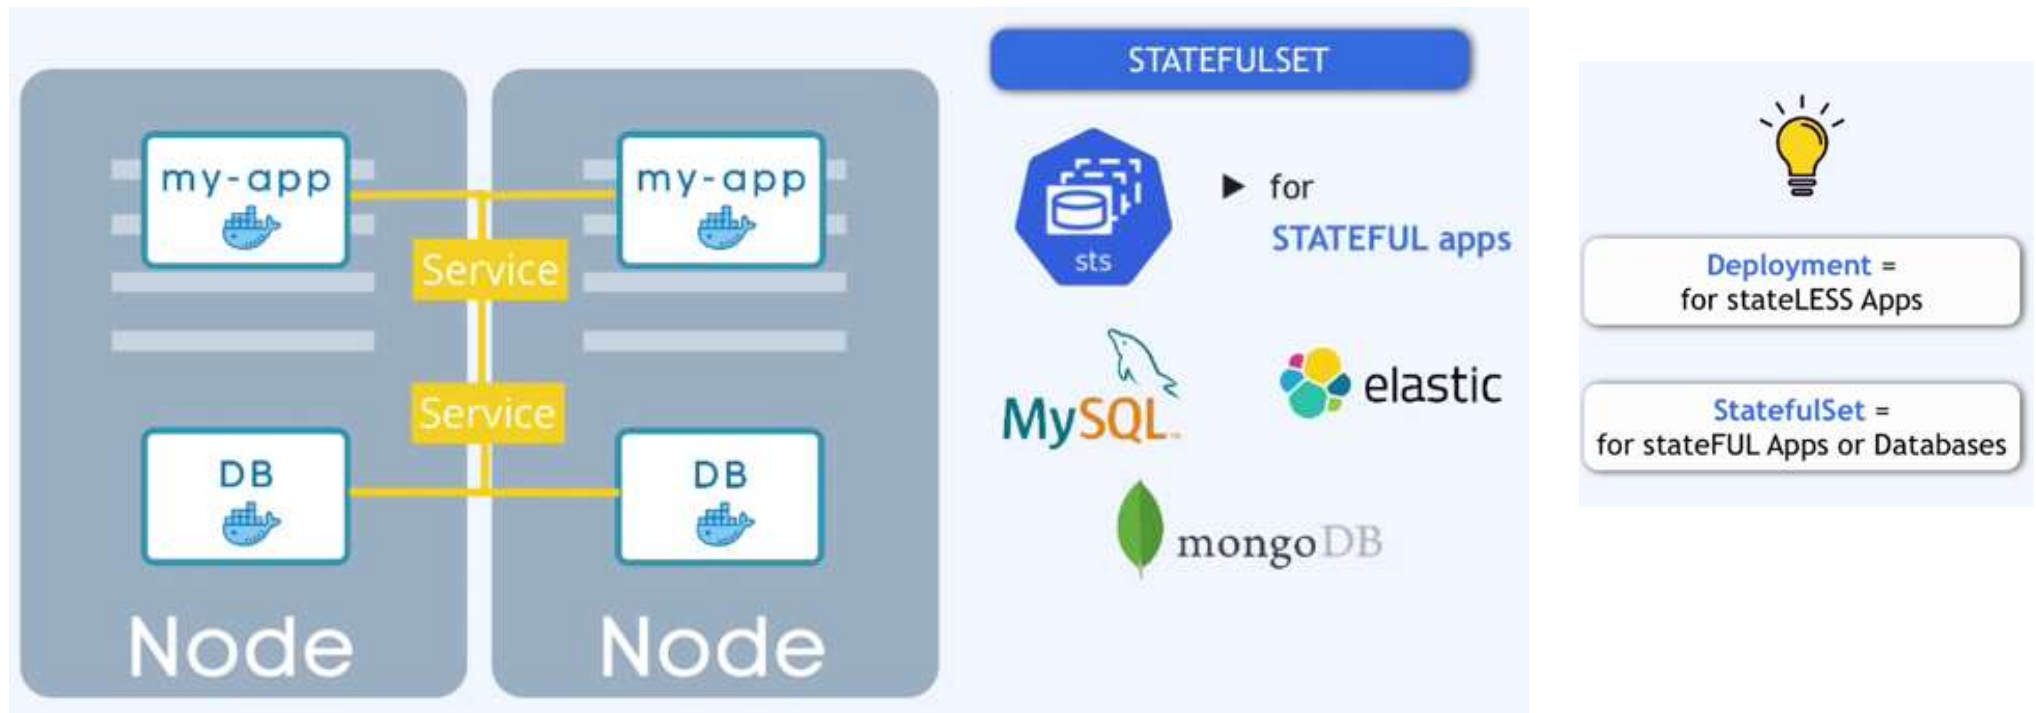
Kubernetes does not manage any persistence. You as Kubernetes user are responsible for managing the data.

Deployment



It's a common practice to keep replicas in separate nodes
Databases cannot be replicated because it has state

Stateful Set



Databases are generally hosted outside the Kubernetes cluster as creating stateful set is not easy

Summary



Kubernetes Configuration

Deployment = template
for creating pods

Declarative in nature


Metadata

Specification

Status -> Automatic

Desired state and actual
state, Kubernetes checks
and if there is a mis-match
it self-heals

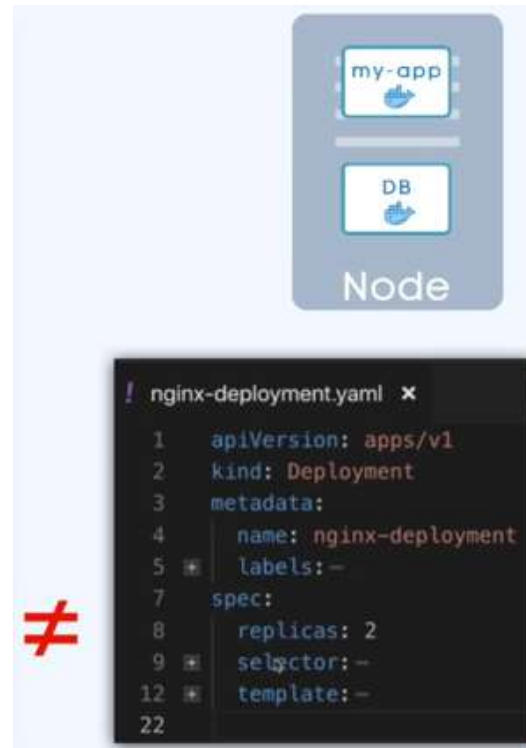
```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: my-app
  labels:
    app: my-app
spec:
  replicas: 2
  selector:
    matchLabels:
      app: my-app
  template:
    metadata:
      labels:
        app: my-app
    spec:
      containers:
        - name: my-app
          image: my-image
          env:
            - name: SOME_ENV
              value: $SOME_ENV
          ports:
            - containerPort: 8080
```



Kubernetes COnfiguration

Kubernetes gets the state of the cluster from etcd

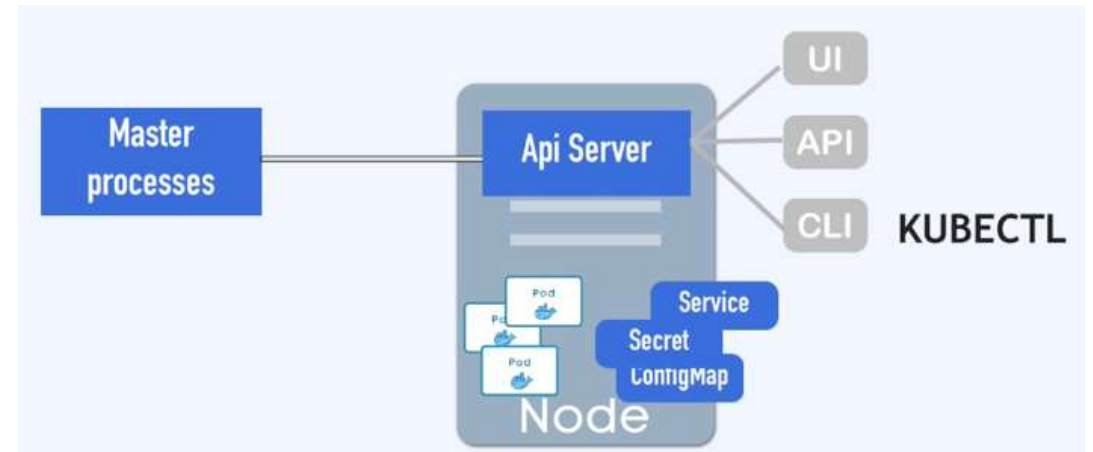
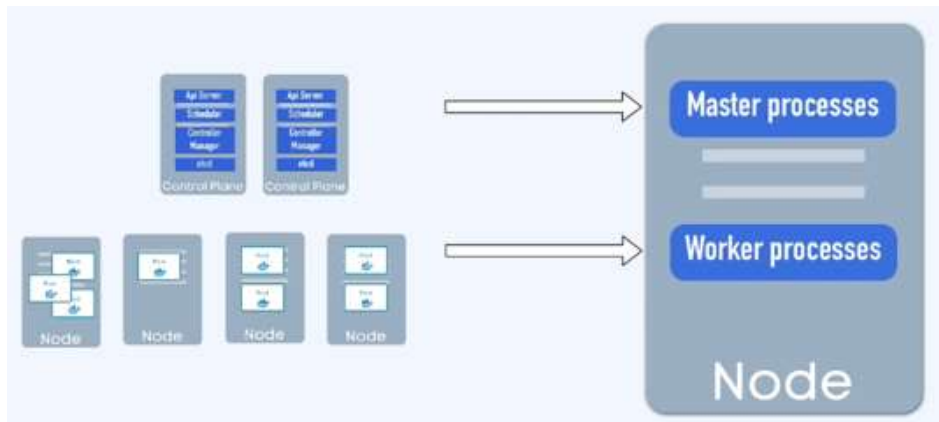
```
status:
  availableReplicas: 1
  conditions:
  - lastTransitionTime: "2020-01-24T10:54:59Z"
    lastUpdateTime: "2020-01-24T10:54:59Z"
    message: Deployment has minimum availability.
    reason: MinimumReplicasAvailable
    status: "True"
    type: Available
  - lastTransitionTime: "2020-01-24T10:54:56Z"
    lastUpdateTime: "2020-01-24T10:54:59Z"
    message: ReplicaSet "nginx-deployment-7d64f4b"
    reason: NewReplicaSetAvailable
    status: "True"
    type: Progressing
  observedGeneration: 1
  readyReplicas: 1
  replicas: 1
  updatedReplicas: 1
```



One more node will be created as the replicas are not matching in the specification

Minikube and Kubectl

Minikube = Local/Test setup for running Kubernetes cluster, master and node processes run in a single machine



Connecting to the node

Can be used not only on the local cluster, but can be used to talk to any type of Kubernetes cluster such as cloud Kubernetes cluster